



Predicting Geolocation from Tweets

Yachay.ai Externship

Overview



01 Model Selection

02 Preprocessing

03 Results

Model Selection





Model Type

KERAS
Functional API



Deep Learning

BERT
Pretrained



Approach

Regression



Tweets

3 NLP
Features

input_2	input:	[(None, 155)]	input_3	input:	[(None, 155)]	input_1	input:	[(None, 3)]
InputLayer	output:	[(None, 155)]	InputLayer	output:	[(None, 155)]	InputLayer	output:	[(None, 3)]

tf_bert_model_1	input:	(None, 155)
TfBertModel	output:	TfBaseModelOutputWithPoolingAndCrossAttentions(last_hidden_state=(None, 155, 768), pooler_output=(None, 768), past_key_values=None, hidden_states=None, attentions=None, cross_attentions=None)

global_average_pooling1d_1	input:	(None, 155, 768)
GlobalAveragePooling1D	output:	(None, 768)

dense	input:	(None, 768)
Dense	output:	(None, 128)

dense_1	input:	(None, 128)
Dense	output:	(None, 64)

dense_2	input:	(None, 64)
Dense	output:	(None, 2)

Preprocessing



NLP FEATURES

Language;
LabelEncoder()

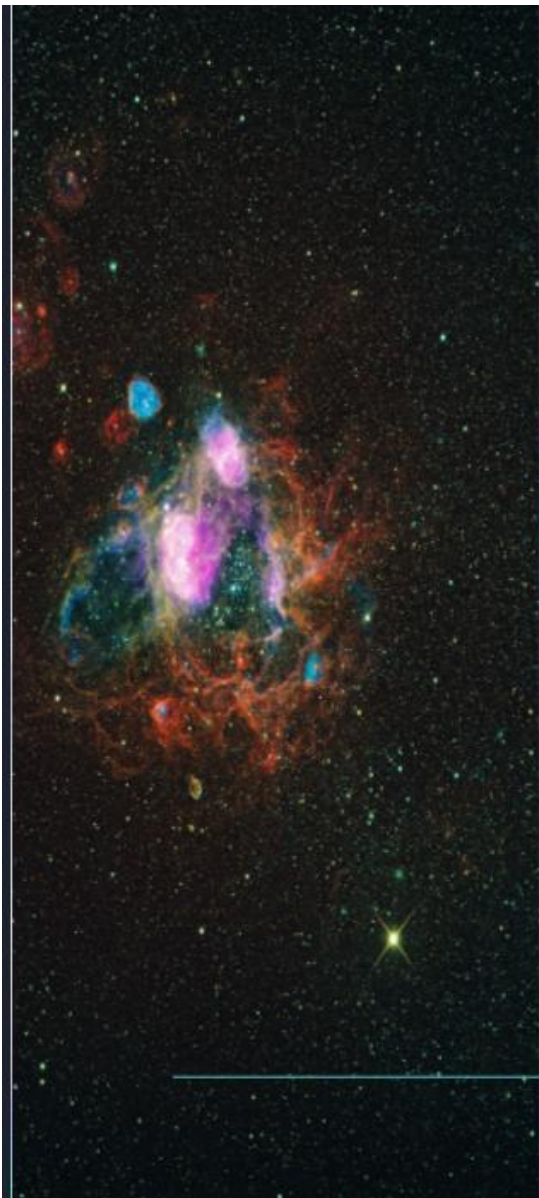
Datetime Features
(dt.month, dt.day)

TWEETS

Tokenized;
AutoTokenizer.from_pretrained
('bert-base-uncased')

Removed N/A





Language Distribution



English Spanish Other

Regions > 10,000 Tweets



Training & Evaluation

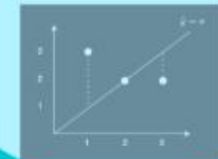
Loss

Haversine Distance



Metric

Mean Squared Error



Learning Rate

0.001

epochs=7

Optimizer

ADAM

Model Structure

```
numerical_input=Input(shape=(num_features,), dtype=tf.float32)
input_ids = Input(shape=(max_seq_length,), dtype=tf.int32)
attention_masks = Input(shape=(max_seq_length,), dtype=tf.int32)

base_model = TFAutoModel.from_pretrained('bert-base-uncased', trainable=False)(input_ids, attention_mask=attention_masks)[0]
x = GlobalAveragePooling1D()(base_model)
concatenated=Concatenate()([numerical_input, x])
x = Dense(128, activation='relu')(x)
x=Dense(64, activation='relu')(x)
output=Dense(2)(x)

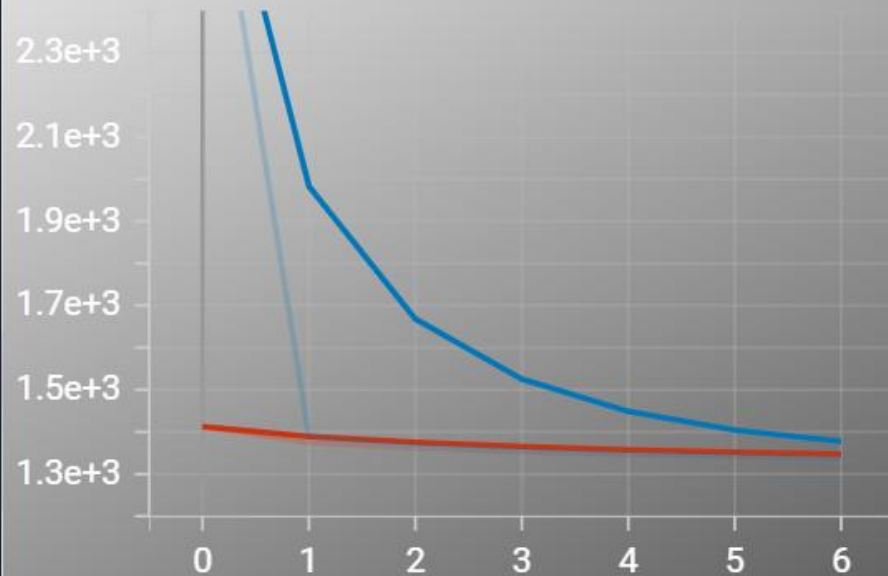
model=Model(inputs=[numerical_input, input_ids, attention_masks], outputs=output)

model.compile(optimizer=optimizer, loss= haversine_distance, metrics=['mse'])

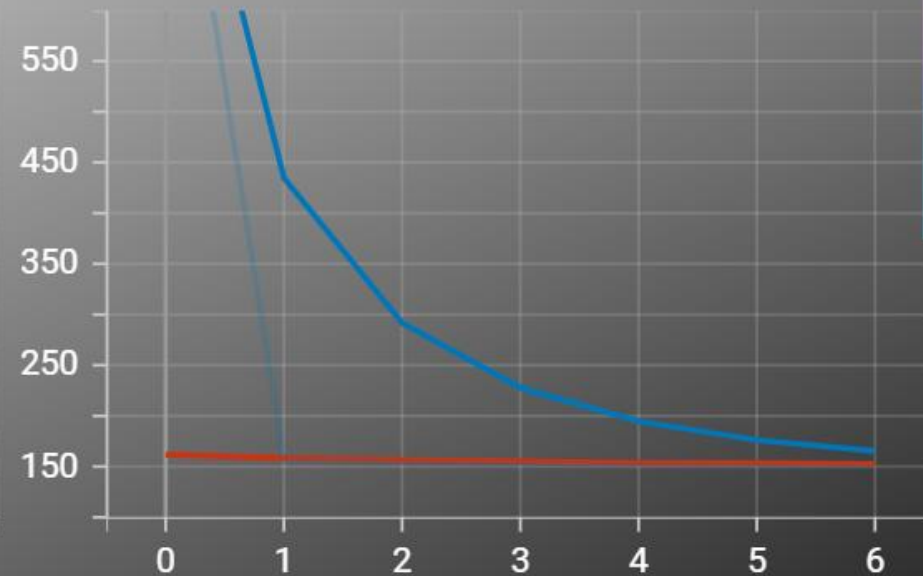
strategy = tf.distribute.OneDeviceStrategy("GPU") if tf.config.list_physical_devices("GPU") else tf.distribute.OneDeviceStrategy("CPU")
history=model.fit([numerical_train_input, train_input_ids, train_attention_masks], y_train,
                  validation_data=([numerical_val_input, val_input_ids, val_attention_masks], y_val),
                  epochs=7, batch_size=2000, callbacks=[lr_scheduler, checkpoint, early_stopping, tensorboard, tensorboard_callback])
```

Results

Loss vs. Epoch



MSE vs. Epoch



— TRAINING — VALIDATION

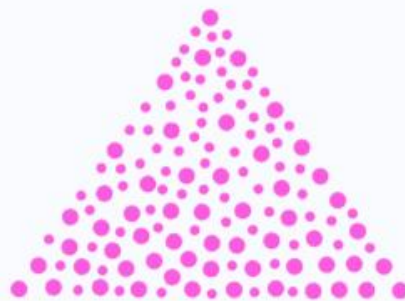
Results

DATA	LOSS	MSE
TRAINING	1356 km	153 km ²
VALIDATION	1345 km	152 km ²
TEST	1334 km	149 km ²

Final Reflection

Limitations

- ☐ Validation data outperformed
- ☐ Short Training



Strengths

- ☐ Utilizes BERT
- ☐ NLP Features
- ☐ Convenient & Buildable



Thank You